

Tutorial: Deep Learning (with Tensorflow)

Ole Salscheider

08.06.2018

Overview

1 Introduction

2 Convolutional Neural Networks

3 Tensorflow


4 Hands-on tutorial

Introduction

History of Artificial Neural Networks

1940s	First ideas, Hebbian rule
1950s–1960s	First successful networks
1970s	Back-propagation algorithm
2010s	Deep neural networks


Artificial neuron




$$o = g \left(\sum_{i=1}^n w_i \cdot x_i + b \right) = g \left(\mathbf{w}^T \mathbf{x} + b \right)$$

$$\mathbf{o} = g(\mathbf{W}\mathbf{x} + \mathbf{b})$$

Artificial neural network




Activation function: tanh (or sigmoid)



- $g(x) = \tanh(x)$
- Was popular in the past...

Activation function: ReLU



- $g(x) = \max(0, x)$
- Converges faster than tanh or sigmoid
- Cheap to compute

Universal approximation theorem

A feed-forward network with a single hidden layer can approximate arbitrary continuous functions (under mild assumptions on the activation function).

A feed-forward network with two hidden layers can approximate arbitrary functions.

But how many neurons do we need? Can we find the weights?

Training

- Training: Minimise cost function $C \Rightarrow$ Learn weights and biases
- Stochastic gradient descent (with batch of size m)

$$w_k \rightarrow w_k - \frac{\eta}{m} \sum_{j=1}^m \frac{\partial C_j}{\partial w_k}$$

$$b_l \rightarrow b_l - \frac{\eta}{m} \sum_{j=1}^m \frac{\partial C_j}{\partial b_l}$$

Back-propagation

- Efficient calculation of $\frac{\partial C}{\partial w}$ and $\frac{\partial C}{\partial b}$
- Consists of forward pass and backward pass
 - ▶ *Forward pass*: Inference, remember weighted inputs and outputs at each layer
 - ▶ *Backward pass*: Propagate the error from the output layer to the earlier layers

Cost function

- Regression

- ▶ L2 loss: $C = \sum_i (y_i - \hat{y}_i)^2$

- Classification (with one-hot vector, e. g. $(0 \ 0 \ 1 \ 0)$)

- ▶ Cross entropy loss: $C = \sum_i -y_i \ln(\hat{y}_i) - (1 - y_i) \ln(1 - \hat{y}_i)$

Weight initialisation

- From pre-trained networks

- Gaussian noise

- ▶ MSRA (for ReLU): $\sigma = \frac{2}{\sqrt{X \cdot Y \cdot (C_{in} + C_{out})}}$
- ▶ Keeps variance of input and output the same

Regularisation


- L2 regularisation
 - ▶ Add squared weights as penalty to the cost function
 - ⇒ Prevents weights from becoming too high
- Dropout
 - ▶ Drop nodes randomly

Deep learning - What is it?

- Deep learning \equiv multiple hidden layers
- Became very popular during the last few years
 - ▶ CNNs
 - ▶ Efficient GPU implementations


Convolutional Neural Networks

Convolution layer




Weights are shared \Rightarrow Prevents the curse of dimensionality

Convolution layer




Weights are shared \Rightarrow Prevents the curse of dimensionality

Convolution layer




Weights are shared \Rightarrow Prevents the curse of dimensionality

Pooling




- Reduces resolution
⇒ computationally cheaper, larger receptive field
- Adds shift invariance

Pooling




- Reduces resolution
⇒ computationally cheaper, larger receptive field
- Adds shift invariance

Pooling




- Reduces resolution
⇒ computationally cheaper, larger receptive field
- Adds shift invariance

Deconvolution layer




- Transposed convolution with stride > 1
- Used for upsampling

Deconvolution layer



- Transposed convolution with stride > 1
- Used for upsampling

Deconvolution layer



- Transposed convolution with stride > 1
- Used for upsampling


ResNet

- Residual neural network
 - Has skip connections / shortcuts
- ⇒ Avoids the *vanishing gradient problem*
- ⇒ Allows to train networks with hundreds to thousands of layers


ResNet modules

- Computes $g(x) = x + f(x)$
- Gives the name to the module

Variant 1



Variant 2



Simple ResNet

- *Very simple ResNet*
 - ▶ We will implement this in the hands-on tutorial
- Usually, ResNets are much deeper and solve more complicated problems
- But this is fast to train on the CPU...



Tensorflow

Tensorflow - Overview

- Developed by Google
 - Published as Open Source in November 2015
 - Currently one of the most widely used deep learning frameworks
-
- Core implemented in C++
 - Training code is (usually) written in Python
 - Supports CPUs, (Nvidia) GPUS and Google TPUs

Tensorflow - Execution modes

- Eager execution
 - ▶ New...
 - ▶ Imperative
 - ▶ Evaluates operations immediately
 - ⇒ Easy to debug
 - ⇒ Natural control flow

Tensorflow - Execution modes

- Graph Execution

- ▶ Python code defines a graph
- ▶ Later, graph is evaluated repeatedly with different input data
- ⇒ Loops have to be defined in graph
- ⇒ Graph can be optimized
- ⇒ Avoids constant switches between Python and C++
- ⇒ Faster

Hands-on tutorial

Online resources

- <http://neuralnetworksanddeeplearning.com>
Good introduction to neural networks
- <https://github.com/ChristosChristofidis/awesome-deep-learning>
Huge collection of links
- <https://github.com/kjw0612/awesome-deep-vision>
Collection of recent papers that cover computer vision applications using deep learning
- <https://www.tensorflow.org/tutorials/>
Tutorials for Tensorflow

Hands-on tutorial: Traffic sign classifier with Tensorflow

Thank you for your kind attention!

Questions?